

Continuation Value is All You Need

“Drop-In” Deep Learning for HA Models with Aggregate Shocks

Jeffrey Sun, University of Toronto

Federal Reserve Board, April 30, 2025

Outline

1. Outline of Algorithm
2. Literature
3. Implementation Details
4. Evaluation
5. Optimizations

Algorithm Outline

Algorithm Outline

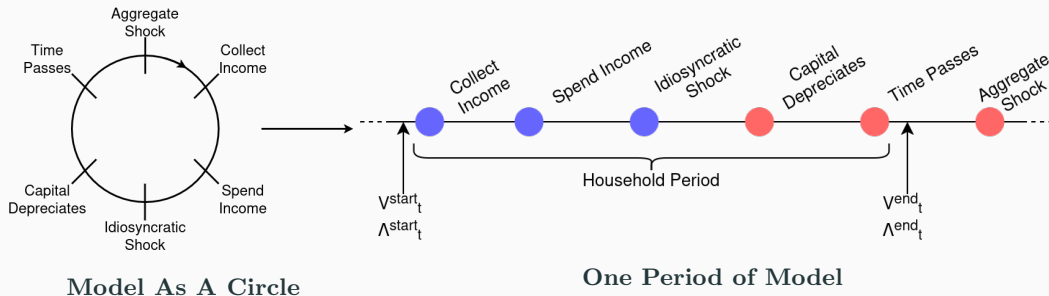
- i. Describe a general class of models
- ii. Make some observations about it
- iii. Outline an algorithm

Model

- Time is discrete
- Heterogeneous agents (Households, Firms, etc.)
- Infinite horizon, defined recursively
- Aggregate shock $Z_t \sim \mathcal{D}()$ hits once per period
- Aggregates (including distributions) are deterministic conditional on $\{Z_t\}_t$

Timing

Denote the spell between aggregate shocks the “household period,”



Denote the value function and aggregate state at start and end of household period,

$$V_t^{\text{start}}, V_t^{\text{end}} \quad \text{and} \quad \Lambda_t^{\text{start}}, \Lambda_t^{\text{end}}.$$

Observation 1: “Within-Period Problem” Solvable Conventionally

Observation 1: There exists a function, typically implementable by conventional methods,

$$\Phi : (\Lambda_t^{\text{start}}, V_t^{\text{end}}) \mapsto (\Lambda_t^{\text{end}}, V_t^{\text{start}}),$$

- The “Within-Period Problem (WPP)” is the set of equations defining Φ .
- Φ is the “Solution” to the WPP.
- A specific implementation of Φ is a “Solver” for the WPP.

Argument:

- By assumption, WPP is aggregate-deterministic.
- Solved like one-period model where V_t^{end} is terminal payoff.
- Often solved by iterating V_t^{end} backward, then simulating Λ_t^{start} forward.

Observation 2: “Between-Period Problem”

Observation 2: There exists a function,

$$\Pi : \Lambda_t^{\text{start}} \mapsto V_t^{\text{end}},$$

as long as the model is Markov and has a solution.

- However, Π can be **infeasible** to implement, especially with aggregate uncertainty.
- Brute force method: VFI over every possible $V_t^{\text{start}}(x_{it}, \Lambda^{\text{start}})$.
- But aggregate state space can be vast, curse of dimensionality binds.

However, this is the *only place* the curse of dimensionality applies...

Observation 3: Π Is Enough For Simulation

Observation 3: If you *somehow* had,

$$\Pi : \Lambda_t^{\text{start}} \mapsto V_t^{\text{end}},$$

then you could simulate forward:

$$\begin{aligned}\Lambda_t^{\text{end}} &= \Phi_1(\Lambda_t^{\text{start}}, \Pi(\Lambda_t^{\text{start}})) \\ \Lambda_{t+1}^{\text{start}} &= \Omega(\Lambda_t^{\text{end}}, Z_t) \\ \Lambda_{t+1}^{\text{end}} &= \Phi_1(\Lambda_{t+1}^{\text{start}}, \Pi(\Lambda_{t+1}^{\text{start}})) \\ &\dots\end{aligned}$$

Indeed, conditional on any candidate Π_0 , we have,

$$(\Lambda^{\text{end}}, V^{\text{start}}) = \tilde{\Phi}(\Lambda^{\text{start}}; \Pi_0) \equiv \Phi(\Lambda^{\text{start}}, \Pi_0(\Lambda^{\text{start}})),$$

and a **computable** law of motion,

$$\Gamma(\Lambda^{\text{start}}; \Pi_0, Z) \equiv \Omega(\tilde{\Phi}_1(\Lambda^{\text{start}}; \Pi_0), Z).$$

Observation 4: Π Recursively Defined

Observation 4: $\Pi : \Lambda^{\text{start}} \mapsto V^{\text{end}}$ can be recursively defined as:

$$\begin{aligned}
 & \forall \Lambda^{\text{start}} \in (\text{Aggregate State Space}) \\
 & V^{\text{end}} \equiv \Pi(\Lambda^{\text{start}}) \\
 & \quad = \mathbb{E}_Z \left[V^{\text{start}'} \mid \Lambda^{\text{start}} \right] \\
 \text{where } & V^{\text{start}'} \equiv \tilde{\Phi}_2 \left(\Lambda^{\text{start}'}; \Pi \right) \\
 & \Lambda^{\text{start}'} \equiv \Gamma(\Lambda^{\text{start}}; \Pi, Z)
 \end{aligned}$$

(i.e. “Projecting 1 step forward” = “Projecting 2 steps forward and solving 1 step back”.)

Essentially a rearrangement of the Bellman equation to define Π .

Π As A Fixed Point

Π is thus a fixed point of the following “lookahead” operator:

$$L\Pi_0(\Lambda^{\text{start}}) \equiv \mathbb{E}_Z \left[\tilde{\Phi}_2(\Gamma(\Lambda^{\text{start}}; \Pi_0, Z); \Pi_0) \right].$$

It remains only to train a neural network Π_0 to minimize:

$$\| \Pi_0(\Lambda^{\text{start}}) - L\Pi_0(\Lambda^{\text{start}}) \|.$$

When Π is a neural network, this is the “Neural Bellman Equation.”

Neural Bellman Equation

Explicitly, for a neural network \mathcal{N}_θ with parameters θ ,

$$\text{NBE}(\theta; \Lambda^{\text{start}}) = \mathcal{N}_\theta(\Lambda^{\text{start}}) - \mathbb{E}_Z \left[\tilde{\Phi}_2(\Gamma(\Lambda^{\text{start}}; \mathcal{N}_\theta, Z); \mathcal{N}_\theta) \right]$$

Measures the failure of \mathcal{N}_θ to produce a lookahead-consistent V^{end} projection at Λ^{start} .

Algorithm Sketch

Let \mathcal{N}_θ be a neural network with parameters θ .

1. Initialize θ somehow.
2. Draw N samples Λ_i^{start} somehow.
3. For each i :

- i. Compute $\widehat{V}^{\text{end}}_i \leftarrow \Gamma \mathcal{N}_\theta(\Lambda_i^{\text{start}})$
- ii. Compute,

$$g_i \leftarrow \frac{\partial}{\partial \theta} \left\| \mathcal{N}_\theta(\Lambda_i^{\text{start}}) - \widehat{V}^{\text{end}}_i \right\|^2$$

4. Update $\theta \leftarrow \theta - \alpha \frac{1}{N} \sum_i g_i$
5. Repeat from Step 2 until convergence.

Literature

Key Departure

Key departure from literature on ML for HA models with aggregate shocks:

- Use **conventional** methods to solve for **everything** but V^{end} : policy, prices, all intermediate value functions, evolution of state, etc.
- In particular, use conventional methods to solve the **policy** function

Key Choices (Within Econ. Lit on ML for HA Models w/ Agg. Shocks)

- Time:
 - **Discrete**
 - Continuous: Gu et al. (2024), Fernández-Villaverde et al. (2023), etc.
- Solution scope
 - **Global:** Most DL based methods
 - Local: Most projection/perturbation methods: Bhandari et al. (2023), Bilal (2023), Auclert et al. (2021), Winberry (2018), etc.
- Policy function
 - **Conventional:** Krusell and Smith (1998), Hull (2015), etc.
 - Deep Learning: Han et al. (2024), Azinovic et al. (2022), Maliar et al. (2021), etc.

Bridging ML and Economics Literatures

One perspective:

- Method is Krusell-Smith, replacing their “finite-moments value function” with a “neural net value function”

Another perspective:

- Method is Approximate Dynamic Programming (ADP) applied to heterogeneous-agent equilibrium models

Implementation

Model Setup

- Standard Krusell-Smith Model [Details](#)
- 65 wealth gridpoints, 3 income gridpoints = 195 idiosyncratic gridpoints

Algorithm Sketch

Let \mathcal{N}_θ be a neural network with parameters θ . \longleftarrow We Are Here

1. Initialize θ somehow.
2. Draw N samples Λ_i^{start} somehow.
3. Compute $\widehat{V}_{\text{end}_i} \longleftarrow \Gamma \mathcal{N}_\theta(\Lambda_i^{\text{start}})$.
4. For each i , compute,

$$g_i \longleftarrow \frac{\partial}{\partial \theta} \left\| \mathcal{N}_\theta(\Lambda_i^{\text{start}}) - \widehat{V}_{\text{end}_i} \right\|^2$$

5. Update $\theta \longleftarrow \theta - \alpha \frac{1}{N} \sum_i g_i$
6. Repeat from Step 2 until convergence.

Basic Neural Network

Idea: Compose matrix multiplication with elementwise nonlinear function.

This can create arbitrary nonlinear relationships between inputs and outputs.

Parameters: $\theta = (M^1, b^1, M^2, b^2)$

Input: x^1

$$y^1 \leftarrow M^1 x^1 + b^1$$

$$x^2 \leftarrow \text{elu}(y^1)$$

$$y^2 \leftarrow M^2 x^2 + b^2$$

$$y^3 \leftarrow \text{elu}(y^2)$$

Output: y^3

Where

$$\text{elu}(x)_i \equiv \begin{cases} x_i & x_i > 0 \\ e^{x_i} - 1 & x_i \leq 0 \end{cases}$$

is an **elementwise** nonlinear function
("activation function.")

Automatic Differentiation

The real magic of neural networks: automatic differentiation.

For a “label” ℓ_x associated with each input x , and a neural network $\mathcal{N}_\theta(x)$, the gradient

$$\frac{\partial}{\partial \theta} \|\mathcal{N}_\theta(x), \ell_x\|^2$$

can be quickly and automatically computed.

- Can often thus find θ such that $\mathcal{N}_\theta(x)$ approximates/predicts ℓ_x well.
- Theorem: a sufficiently large neural net can approximate *any* θ arbitrarily well

Specific Neural Net Implementation: Step 1

Step 1: Use “Generalized Moments” mean field game-inspired method of Han, Yang, and E (2025) to summarize household state distribution μ .

$$\forall x \in X :$$

$$\gamma_x^{GM,1} \leftarrow \text{elu} \left(M_{10 \times 2}^1 \begin{pmatrix} k_x \\ z_x \end{pmatrix} + b^{GM,1} \right)$$

$$\gamma_x^{GM,2} \leftarrow \text{elu} (M_{10 \times 10}^2 \gamma_x^{GM,1} + b^{GM,2})$$

$$\gamma^{GM,3} \leftarrow \sum_{x \in X} \mu(x) \gamma_x^{GM,2}$$

$$\gamma^{GM} \leftarrow \text{elu} (M_{10 \times 10}^3 \gamma^{GM,3} + b^{GM,3})$$

Output: γ^{GM} , containing information about aggregate state distribution.

Specific Neural Net Implementation: Step 2

Step 2: Combine state distribution summary γ^{GM} with idiosyncratic state x and aggregate productivity A .

$$\gamma_x^{HH,1} \leftarrow \text{elu} \left(M_{10 \times 2}^{HH,1} \begin{pmatrix} k_x \\ z_x \end{pmatrix} + b^{HH,1} \right)$$

$$\gamma^{A,1} \leftarrow \text{elu} \left(M_{10 \times 1}^{A,1} A + b^{A,1} \right)$$

$$\gamma^A \leftarrow \text{elu} \left(M_{10 \times 10}^{A,2} \gamma^{A,1} + b^{A,2} \right)$$

$$\gamma_x^{HH,2} \leftarrow \gamma_x^{HH,1} + \gamma^{GM} + \gamma^A$$

Output: $\gamma_x^{HH,2}$, containing information about idiosyncratic x and aggregate $\Lambda^{\text{start}} = (\mu, A)$.

Specific Neural Net Implementation: Step 3

Step 3: Put $\gamma_x^{HH,2}$ through three more layers:

$$\gamma^{HH,3} \leftarrow \text{elu}\left(M_{8 \times 10}^{HH,3} \gamma_x^{HH,2} + b^{HH,3}\right)$$

$$\gamma^{HH,4} \leftarrow \text{elu}\left(M_{8 \times 8}^{HH,4} \gamma^{HH,3} + b^{HH,4}\right)$$

$$\gamma^{HH,5} \leftarrow \text{elu}\left(M_{5 \times 8}^{HH,5} \gamma^{HH,4} + b^{HH,5}\right)$$

$$\widehat{V}(x; \Lambda^{\text{start}}, \theta) \equiv M_{1 \times 5}^{HH,6} \gamma^{HH,5} + b^{HH,6}$$

Output: $\widehat{V}(x; \Lambda^{\text{start}}, \theta)$, an estimate of the state-contingent value $V(x; \Lambda^{\text{start}})$.

Krusell-Smith Comparison

The Krusell-Smith *method* is essentially the same, but with \hat{V} given by,

$$\begin{aligned}\bar{k} &\leftarrow \sum_{x \in X} \mu(x) k_x \\ \gamma^{HH} &\leftarrow \begin{pmatrix} v(k, 1; \bar{k}, z_g) & v(k, 1; \bar{k}, z_b) & v(k, 0; \bar{k}, z_g) & v(k, 0; \bar{k}, z_b) \end{pmatrix}' \\ \hat{V}^{KS}(x; A, \mu) &\equiv \begin{pmatrix} \mathbb{1}_{[A=1, z_x=z_g]} & \mathbb{1}_{[A=1, z_x=z_b]} & \mathbb{1}_{[A=0, z_x=z_g]} & \mathbb{1}_{[A=0, z_x=z_b]} \end{pmatrix} \gamma^{HH}\end{aligned}$$

No need to interpret as “predicting the law of motion of aggregate capital.” It’s just value function approximation!

Data Representation

- Represent $V_t^{\text{start}}, \mu_t^{\text{start}}$ by data arrays $A_t^{V^{\text{start}}}, A_t^{\mu^{\text{start}}}$. Similarly, $A_t^{V^{\text{end}}}, A_t^{\mu^{\text{end}}}$.
- Φ is then implemented as a function taking arrays to arrays:

$$\Phi : \left(A^{V^{\text{end}}}, A^{\Lambda^{\text{start}}}, S_t^{\text{start}} \right) \mapsto \left(A^{V^{\text{start}}}, A^{\Lambda^{\text{end}}}, S_t^{\text{end}} \right)$$

where S_t^{start} and S_t^{end} are other state variables.

- Implemented **conventionally** (no neural net)

Algorithm Sketch

Let \mathcal{N}_θ be a neural network with parameters θ .

1. Initialize θ somehow. \longleftarrow We Are Here
2. Draw N samples Λ_i^{start} somehow.
3. Compute $\widehat{V}_{\text{end}_i} \longleftarrow \Gamma \mathcal{N}_\theta(\Lambda_i^{\text{start}})$.
4. For each i , compute,

$$g_i \longleftarrow \frac{\partial}{\partial \theta} \left\| \mathcal{N}_\theta(\Lambda_i^{\text{start}}) - \widehat{V}_{\text{end}_i} \right\|^2$$

5. Update $\theta \longleftarrow \theta - \alpha \frac{1}{N} \sum_i g_i$
6. Repeat from Step 2 until convergence.

Initializing θ

Starting with a completely random θ may lead to terrible initial simulations.

Several options here:

1. Pre-train θ to minimize

$$\left\| \mathcal{N}_{\theta}(\Lambda^{\text{start}})(x) - \left(k_x + \frac{z_x}{1 - \beta} \right) \right\|$$

2. Pre-train θ on $(\Lambda^{\text{start}}, V^{\text{end}})$ data-label pairs from:

- 2.1 Deterministic version of model

- 2.2 Krusell-Smith *method* solution of full model

Algorithm Sketch

Let \mathcal{N}_θ be a neural network with parameters θ .

1. Initialize θ somehow.
2. Draw N samples Λ_i^{start} somehow. \longleftarrow We Are Here
3. Compute $\widehat{V_i^{\text{end}}} \longleftarrow \Gamma \mathcal{N}_\theta(\Lambda_i^{\text{start}})$.
4. For each i , compute,

$$g_i \longleftarrow \frac{\partial}{\partial \theta} \left\| \mathcal{N}_\theta(\Lambda_i^{\text{start}}) - \widehat{V_i^{\text{end}}} \right\|^2$$

5. Update $\theta \longleftarrow \theta - \alpha \frac{1}{N} \sum_i g_i$
6. Repeat from Step 2 until convergence.

Sampling Λ^{start}

Options:

1. Draw from simple (e.g. uniform) distribution
2. Draw from ergodic distribution induced by θ
3. Draw from ergodic distribution induced by some other Π_0 , in MCMC spirit
 - E.g. Krusell-Smith
4. Take current data state as starting point, sample from possible states within T periods of present

Sampling Λ^{start}

I use MCMC-style sampling:

1. Simulate forward T “burn-in” periods using

$$\Lambda_{t+1}^{\text{start}} \longleftarrow \Gamma(\Lambda_t^{\text{start}}; \mathcal{N}_\theta, Z)$$

2. For $t > T$, add Λ_t^{start} to sample if $t \equiv 0 \pmod{n}$ for some n

Basic challenge: Data generation depends on \mathcal{N}_θ .

Algorithm Sketch

Let \mathcal{N}_θ be a neural network with parameters θ .

1. Initialize θ somehow.
2. Draw N samples Λ_i^{start} somehow.
3. Compute $\widehat{V_i^{\text{end}}} \leftarrow \Gamma \mathcal{N}_\theta(\Lambda_i^{\text{start}})$. **← We Are Here**
4. For each i , compute,

$$g_i \leftarrow \frac{\partial}{\partial \theta} \left\| \mathcal{N}_\theta(\Lambda_i^{\text{start}}) - \widehat{V_i^{\text{end}}} \right\|^2$$

5. Update $\theta \leftarrow \theta - \alpha \frac{1}{N} \sum_i g_i$
6. Repeat from Step 2 until convergence.

Within-Period Problem

Within-period problem: find,

$$\Phi : (\Lambda_t^{\text{start}}, V_t^{\text{end}}) \mapsto (\Lambda_t^{\text{end}}, V_t^{\text{start}}).$$

Performance strategy: break into “stages”: effectively “sub-periods.”

Krusell-Smith model is easy, but we can get much richer.

Sub-Periods: “Conventionally” Solved but *Modular*



Similar decomposition to Auclert et al. (2023)

Multiplicative performance benefits:

1. Simple, optimized, prepackaged for CPU/GPU/cluster
2. Search over one dimension at a time
3. Identify and overcome bottlenecks (interpolation, grid search) in non-vectorized Julia

Fixing Bottlenecks: Dynamic Information Sharing

Example: Consumption-saving by grid search

- If $MPC \geq 0$, then my optimal saving is between my wealth-neighbors'
- Don't need to search over entire axis!
- By “sharing” information between wealth-neighbors: $O(N^2) \rightarrow O(N \log N)$
- Incompatible with vectorization (Python, Matlab) but fast in Julia

Pre-Packaged Stages

Another benefit of stages: can be pre-prepared

- Possibly separate project
- Aim to provide highly efficient state implementations for CPU, GPU, and cluster

Household Problem Code Example

```
function solve_period!(prealloc, V_next, params)
    V_preshock = get_V_preshock(prealloc, V_next)

    V_consume = get_V_preconsume(V_preshock, prealloc)

    V_income = get_V_preincome(V_consume, prealloc, params)
    enforce_borrowing_constraint!(V_preincome, prealloc)

    V_remove = get_V_remove(V_preincome, prealloc, params)

    V_end = YOUR_CODE_HERE(V_remove, prealloc, params)
end
```

Solving for Prices

Options:

1. Solve analytically (e.g. Krusell-Smith)
2. Solve via rootfinding (accurate but slow)
3. Use additional neural network (e.g. Azinovic et al. (2022), Deep Equilibrium Nets)

Within-Period Problem


Challenge: Performance of Within-Period Solver is crucial.

Accuracy


Accuracy

- Strength: $\Phi : (\Lambda_t^{\text{start}}, V_t^{\text{end}}) \mapsto (\Lambda_t^{\text{end}}, V_t^{\text{start}})$ is immediately correct from the beginning
- Only need to check Neural Bellman Equation error
- Weakness: Computing Euler Equation errors nontrivial
- Hardware: One laptop CPU thread (i9-13900H)

Learning Curve Comparison



V error as share of V variance by training epoch



Relative V error by training time (CPU)

- Both stop improving within about 300 epochs (122s for NN, 13s for KS)
- After Epoch 300, mean \log_{10} error is -4.431 for NN, -3.152 for KS

Optimizations

Benchmark

- 1000 locations, 129 wealth states, 5 income types, 6 age groups = 3.87m gridpoints
- Single-thread CPU
- Language: Julia
- Strawman: Jeffrey, May 2023
- One evaluation of household problem
- Initial time: 218s (3m38s)

Low-Hanging Fruit

- **Initial:** 218s
- **Memory Preallocation:** 218s \rightarrow 152s
- **(Almost) Automatic Multithreading:** 152s \rightarrow 49s
- **32 Bit Precision:** 49s \rightarrow 31.9s

Individual Stage Optimizations

Choose Location	25.2s \rightarrow 9.78s \rightarrow 0.053s
Receive Income	0.38s \rightarrow 0.019s
Choose Consumption	0.498s \rightarrow 0.025s
Income Shock	4.52s \rightarrow 0.028s

Overall: 31.9s \rightarrow 0.353s (= 0.126s listed stages + 0.228s other)

Choose Location

Let $V_{ts\iota}^{\text{start}}(\ell) = V_{ts}^{\text{start}}(k_{t-1}, z_{t\iota}, \ell, a_{t\iota})$, $V_{ts\iota}^{\text{end}}(\ell)$ similar

where ι indexes all household types up to location.

The i.i.d. Gumbel location preference shocks imply:

$$\exp(\psi V_{ts\iota}^{\text{start}}(\ell)) = \sum_{\ell'} \exp(\psi(V_{ts\iota}^{\text{end}}(\ell') - D_{\ell\ell'}))$$

$$P(\ell' = \ell_0 \mid \ell) = \frac{\exp(\psi(V_{ts\iota}^{\text{end}}(\ell') - D_{\ell\ell'}))}{\exp(\psi V_{ts\iota}^{\text{start}}(\ell))}$$

$$\lambda_{ts\iota}^{\text{end}}(\ell) = \sum_{\ell'} P(\ell' = \ell \mid \ell) \lambda_{ts\iota}^{\text{start}}(\ell')$$

First optimization: Precompute $\exp(\psi V_{ts\iota}^{\text{start}}(\ell))$, then $P(\ell' = \ell_0 \mid \ell)$, then $\lambda_{ts\iota}^{\text{end}}(\ell)$

Time: 25.2s \rightarrow 9.78s

Choose Location

Second optimization: observe that (with \otimes and \oslash elementwise mult. and div.)

$$\tilde{V}_{ts}^{\text{start}} = D \tilde{V}_{ts}^{\text{end}}$$

$$\Lambda_{ts}^{\text{end}} = \tilde{V}_{ts}^{\text{end}} \otimes (D' \Lambda_{ts}^{\text{start}} \oslash \tilde{V}_{ts}^{\text{start}})$$

where matrices $D_{\ell\ell'} = \exp(-\psi D_{\ell\ell'})$

$$\left(\tilde{V}_{ts}^{\text{start}}\right)_{\ell\iota} = \exp\left(\psi V_{ts\iota}^{\text{start}}(\ell)\right)$$

$$\left(\tilde{V}_{ts}^{\text{end}}\right)_{\ell\iota} = \exp\left(\psi V_{ts\iota}^{\text{end}}(\ell)\right)$$

$$\left(\Lambda_{ts}^{\text{end}}\right)_{\ell\iota} = \lambda_{ts\iota}^{\text{end}}(\ell)$$

No matter the size of the state space, just two matrix multiplications!

Time: 9.78s \rightarrow 0.053s

The Power of Matrix Multiplication

Why is matrix multiplication 200 faster than an explicit loop?

- Surprising algorithms exist to multiply two matrices in as little as $O(n^{2.371552})$ time
- Most CPUs have specialized hardware for matrix multiplication
- Pretty much exactly the same thing works for CES production functions, etc.
- Similar approach to optimizing income shocks, or any finite-state Markov process

Choose Consumption

- Strategy: Gridsearch
- If $MPC \geq 0$, then my optimal saving is between my wealth-neighbors'
- Don't need to search over entire axis!
- By “sharing” information between wealth-neighbors: $O(N^2) \rightarrow O(N \log N)$
- Incompatible with vectorization (Python, Matlab) but fast in Julia
- Similar approach for linear interpolation of many gridpoints

Global Solution

- A neural network is trained to predict V^{end} . Everything else is conventional
- In particular, no neural network used to approximate policy function

Conclusion

Conclusion

- Describe a solution method for HA models with agg. shocks that overcomes curse of dimensionality (of aggregate state)
- Method uses neural nets only where needed – solution otherwise conventional
- Much complexity is offloaded to **Within-Period Problem** solver
- In other work, provide tools to implement WPP flexibly, easily, performantly

Discussion

- Key advantages:
 - Complex household problems supported
 - No need to train policy network
- Disadvantages:
 - Individual state must be low-dimensional (≤ 6 or so)
 - In more complex models, prices require inner loop around IPP or price neural net à la Azinovic et al. (2022)
- Future work:
 - Assess other error metrics, e.g. Euler equation error, risk premia
 - Compare economics of solutions

Appendix

Definition: Aggregate State

Λ_t^{start} (Λ_t^{end}) is the aggregate state at the start (end) of household period t .

$$\Lambda_t^{\text{start}} \equiv (\mu_t^{\text{start}}, S_t^{\text{start}}) \qquad (\Lambda_t^{\text{end}} \equiv (\mu_t^{\text{end}}, S_t^{\text{end}}))$$

where at the start (end) of household period t ,

- μ_t^{start} (μ_t^{end}) is the distribution of idiosyncratic states $x^{\text{start}} \in X^{\text{start}}$ ($x^{\text{end}} \in X^{\text{end}}$).
- S_t^{start} (S_t^{end}) is all other aggregate state.

Without loss of generality, parametrize the aggregate shock so that,

$$\Lambda_{t+1}^{\text{start}} = \Omega(\Lambda_t^{\text{end}}, Z_t) \quad \text{where} \quad Z_t \sim \mathcal{D}().$$

Definition: Value Function

For household i , $V_t^{\text{start}}(x_i)$ is NPV of future utility given state x_i^{start} at the beginning of t ,

$$V_t^{\text{start}}(x_i) \equiv V^{\text{start}}(x_{it}, \Lambda^{\text{start}}) \equiv \mathbb{E} \left[\sum_{s=0}^{\infty} \beta^s u_{i,t+s} \mid x, \Lambda^{\text{start}} \right].$$
$$V_t^{\text{end}}(x_i) \equiv V^{\text{end}}(x_{it}, \Lambda^{\text{start}}) \equiv \mathbb{E} \left[\sum_{s=1}^{\infty} \beta^s u_{i,t+s} \mid x, \Lambda^{\text{start}} \right].$$

- V_t^{start} refers to the whole *function*, e.g. an array on a grid over X^{start} .
- Conditioning on Λ_t^{start} and Λ_t^{end} are equivalent by aggregate-determinism within t .

Algorithm Details

- I use a gridded CDF representation of Λ , but using a finite number of agents is also possible. However, they have to interpolate over continuation V
- Training data for each simulated period is $A_{it}^{V\text{end}}$ together with

$$\mathbb{E}\left[A_{i,t+1}^{V,\text{start}} \mid \Gamma_{it}\right] = \frac{1}{K} \sum_{k=1}^K \text{IPP}_V\left(\mathcal{N}(\Omega(\Gamma_{it}, k) ; \theta_i), A_{it}^{\Lambda\text{end}}, \Omega(\Gamma_{it}, k)\right)$$

- For large models, if memory is constrained, you can update θ_i as you go, accumulating gradients but not storing the entire simulation

Neural Network Implementation

Neural net $\mathcal{N}(A_{\Gamma}^{\Lambda\text{start}}; \theta)$ has following components:

1. Generalized-moment of Han et al. (2023):

$$\text{GM}_{\Gamma} = \sum_{j \in J} (A_{\Gamma}^{\Lambda\text{start}})_j \mathcal{N}_{\text{GM}}(X_j)$$

2. One layer ($1 \Rightarrow 10$) neural net on aggregate productivity A
3. Dense feedforward neural net on input: $(X_j, \text{GM}_{\Gamma}, \mathcal{N}_A(A))$
 - Three hidden layers with 8, 8, and 5 neurons
 - Elu activation

Krusell-Smith Model Details

- 65 wealth gridpoints
- 3 income gridpoints
- $\beta = 0.98$
- Income process by Tauchen discretization with persistence 0.95 and std 0.1
- Log-linear wealth grid from 1k to 10m
- Income states: 15.4k, 40.3k, 105.4k
- Risk aversion: 0.9
- Capital share: 0.36
- Depreciation rate: 0.025

Krusell-Smith Method Details

- 5 aggregate capital gridpoints: 100k, 150k, 200k, 250k, 300k
- Linear extrapolation outside aggregate capital grid
- 2 aggregate productivity states: 0.5 and 1.0
- Unlike Krusell and Smith (1998), use gridded CDF population distribution representation for cleaner comparison

[Back](#)